

CKA Solutions

Monday, October 25, 2021 11:11 AM

Cluster Administration

Task #1

On master node, upgrade Kubernetes to latest version, including kubelet and kubectl.

Solution Steps

1. **kubeadm upgrade plan**
2. **kubectl drain cka-cluster380-master1** (we need to upgrade only master- so we need to drain it. It gives us some warnings and the final command is like this
 - a. **kubectl drain cka-cluster380-master1 --ignore-daemonsets -- delete-local-data**
 - i. Now we delete one more pod via : **kubectl delete pod metrics-server-84645594-h89sg -n kube-system**
 - b. Rerun the drain command after deleting the pod:
kubectl drain cka-cluster380-master1 --ignore-daemonsets -- delete-local-data
3. **apt-get update** (to update the repository)
4. **apt-get install kubelet kubeadm kubectl -y**
5. **systemctl restart kubelet**
6. **kubectl uncordon cka-cluster380-master1**

Screen Command for multiple windows

1. **apt-get install screen**
2. **screen**
3. Run your commands on this screen
4. Use
 - a. **CTRL + A + D** to detach from screen
 - b. **screen -r** to resume to the session
 - c. **CTRL + D** to kill the session
 - d. **screen -ls** to check the available screens

Task #2

One of the Kubernetes nodes is in NotReady state. Investigate the root cause and bring the node to a Ready state.

Note: You can ssh to the appropriate node using: ssh nodeIP

Any changes made should be made permanent

Solution Steps

1. kubectl get nodes -o wide (to check the node which is not ready and its IP address)
2. **ssh root@cka-cluster380-node4 or ssh root@172.16.3.3**
3. **systemctl status docker** (**working**)
4. **systemctl status kubelet** (**not working**)
5. **systemctl start kubelet**
6. **systemctl enable kubelet** (to make it permanent)
7. **systemctl status kubelet** (**working**)

Task #3

Set the node5 as unavailable and reschedule all the pods running on it

Solution Steps

1. **kubectl get nodes** to get the node's full name (cka-cluster380-node5)
2. **kubectl drain cka-cluster380-node5**
3. Update command with the additional flags suggested by terminal:
a. **kubectl drain cka-cluster380-node5 --ignore-daemonsets**
4. **kubectl get nodes** to confirm that the node is in scheduling disabled status

```
root@cka-cluster380-master1:~/metrics-server# kubectl get nodes
NAME           STATUS      ROLES   AGE    VERSION
cka-cluster380-master1  Ready, SchedulingDisabled  master   43h    v1.19.1
cka-cluster380-node1   Ready      <none>  42h    v1.19.1
cka-cluster380-node2   Ready      <none>  42h    v1.19.1
cka-cluster380-node3   Ready      <none>  42h    v1.19.1
cka-cluster380-node4   Ready      <none>  42h    v1.19.1
cka-cluster380-node5   Ready, SchedulingDisabled  <none>  42h    v1.19.1
root@cka-cluster380-master1:~/metrics-server#
```

Pods, Resource Requests/Limits

Task #4

Create a pod named **monolithic** with 2 containers running off of **nginx** + **memcached** images

Solution Steps

1. **kubectl run monolithic--image=nginx --dry-run=client -o yaml >task4.yaml**
2. **vi task4.yaml**
 - a. Add another image" memcached
 - b. wq!
3. **kubectl apply -f task4.yaml**
4. **kubectl get pods**
5. Confirm that the monolithic pod has 2/2 containers running

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~/metrics-server# kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
call           1/1     Running   0          2d8h
jenkins-6f6c59db6c-52wgc  1/1     Running   0          2d8h
jenkins-6f6c59db6c-lqf2r  1/1     Running   0          2d8h
logger          1/1     Running   0          2d8h
marketplace-app 1/1     Running   0          2d8h
mobile-app-575f49855c-vknh6 1/1     Running   0          13h
monolithic      2/2     Running   0          15s
probe-app        1/1     Running   0          2d8h
root@cka-cluster380-master1:~/metrics-server#
```

Task #5

Locate the logs of Pod **marketplace-app** and filter lines corresponding to message **dummy** and send them to **/tasks/JZNEIGGD1/marketplace-app.log**

Solution Steps

1. **kubectl logs marketplace-app |grep dummy > /tasks/JZNEIGGD1/marketplace-app.log**
2. **cat /tasks/JZNEIGGD1/marketplace-app.log**

```
root@cka-cluster380-master1:~# cat /tasks/JZNEIGGD1/marketplace-app.log
::1 - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"
::1 - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"
::1 - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"
::1 - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"
```

```
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
::1 - - "OPTIONS * HTTP/1.0" 200 - "-" "Apache/2.4.25 (Debian) OpenSSL/1.0.2k-fips PHP/5.6.39 (internal dummy connection)"  
root@cka-cluster380-master1:~#
```

Task #6

Update an existing Pod named **logger**:

- Run sidecar container off of busybox image, name it logger
- Shared volume **/var/log** should be mounted on **both containers**, which does not persist when the pod is deleted
- The logger container should run following command: **tail -f /var/log/input.log**

Solution Steps

- We will copy the existing .yaml configuration of logger
 - kubectl get pod logger -o yaml > logger.yaml**
- We will add the second container named logger from busybox image
 - vi logger.yaml**
 - add container (name: logger, image: busybox)
- We will add the arguments to the logger container
 - kubectl run test --image=busybox --dry-run=client -o yaml --tail -f /var/log/input.log**
 - copy and paste the lines about arguments to logger.yaml

```
spec:  
  containers:  
    - image: evolvecybertraining/logger-app  
      imagePullPolicy: Always  
      name: app  
      resources: {}  
      terminationMessagePath: /dev/termination-log  
      terminationMessagePolicy: File  
      volumeMounts:  
  
    - image: busybox  
      name: logger  
      args:  
        - tail  
        - -f  
        - /var/log/input.log
```

4. We will mount an emptyDir volume to both containers based on [Kubernetes resource page](#)
 - a. Add the lines about **volume** to entire pod
 - b. Add the lines about **volume mount** to both containers

```
containers:
- image: evolvecybertraining/logger-app
  imagePullPolicy: Always
  name: app
  resources: {}
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
  volumeMounts:
  - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
    name: default-token-pdc1k
    readOnly: true

  volumeMounts:
  - mountPath: /var/log
    name: cache-volume

- image: busybox
  name: logger
  args:
  - tail
  - -f
  - /var/log/input.log

  volumeMounts:
  - mountPath: /var/log
    name: cache-volume
```

5. Apply the configuration file forcefully
 - a. **kubectl apply --force -f logger.yaml**
6. Confirm the pod is running with two containers.
 - a. **kubectl get pods**

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

root@cka-cluster380-master1:~# kubectl get pods |grep logger
logger   2/2     Running   0        26s
root@cka-cluster380-master1:~# █
```

GREAT!!

Task #7

Create a Pod with 128MB memory and 100M CPU limit

Solution Steps

1. Create a random Pod with dry-run:
 - a. **kubectl run limitpod --image=nginx --dry-run=client -o yaml > limitpod.yaml**
2. Edit the pod definition file to configure the pod definition for limits based on the [Kubernetes documentation page](#) (Search string: **Limit**).
 - a. **vi limitpod.yaml**
3. Create the Pod:

- a. **kubectl apply -f limitpod.yaml**
- 4. Verify the Pod
 - a. **kubectl get pods**

```

TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
call          1/1     Running   0          2d9h
jenkins-6f6c59db6c-52wgc  1/1     Running   0          2d9h
jenkins-6f6c59db6c-lqf2r  1/1     Running   0          2d9h
limitpod      1/1     Running   0          19s
logger         1/1     Running   0          2d9h
marketplace-app  1/1     Running   0          2d9h
mobile-app-575f49855c-vknh6  1/1     Running   0          14h
monolithic    2/2     Running   0          53m
probe-app     1/1     Running   0          2d9h
root@cka-cluster380-master1:~# █

```

Liveness and Readiness Probes

Task #8

A Pod in your cluster often gets locked in a broken state, you have to remedy this situation.

The desired behavior is to restart the Pod when it returns HTTP 500 response and never route Service traffic to the Pod while it is failing. Complete the following:

- The Pod has endpoint **/healthz** that will indicate if the application is still working as expected by returning HTTP 200 response. If the endpoint returns HTTP 500, the application is no longer responsive.
- The Pod has another endpoint **/started**, that will indicate if it can accept traffic by returning HTTP 200 response. If the endpoint returns HTTP 500, the application has not yet finished initialization.
- Configure the probe-app Pod to use these endpoints
- The probes should use **port 8080**
- Do not edit probe-http service

Solution Steps

1. Get the .yaml configuration of the existing probe-app Pod and save it in a .yaml file -
 - a. **kubectl get pod probe-app -o yaml > probe.yaml**
2. Edit the .yaml file and add liveness and readiness probes to the container as indicated in the [K8 resources page](#).

- a. **vi probe.yaml**
3. Create the resource by forcefully applying the changes.
 - a. **kubectl apply --force -f probe.yaml**

```
spec:
  containers:
    - image: evolvecybertraining/probe-app:v1
      imagePullPolicy: Always
      name: probe-app
      resources: {}
      livenessProbe:
        httpGet:
          path: /healthz
          port: 8080
        initialDelaySeconds: 3
        periodSeconds: 3

      readinessProbe:
        httpGet:
          path: /started
          port: 8080
        initialDelaySeconds: 3
        periodSeconds: 3

      terminationMessagePath: /dev/termination-log
      terminationMessagePolicy: File
      volumeMounts:
        - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
          name: default-token-pdclk
          readOnly: true
      dnsPolicy: ClusterFirst
      enableServiceLinks: true
```

```
root@cka-cluster380-master1:~# kubectl apply --force -f task8.yaml
Warning: kubectl apply should be used on resource created by either kubectl create --save-config or kubectl apply
pod/probe-app configured
```

GREAT!

Labels, Taints and Tolerations

Task #9

In production namespace, out of the Pods labeled **group=tools**, find Pod which is using the most CPU resources and pipe the name to /tasks/JSCUERBQ1/pod.txt

Solution Steps

1. Rank the pods which uses the highest CPU
 - a. **kubectl top pods -n production**
2. Check the labels of the pods

- a. `kubectl get pods --show-labels -n production`
3. See the pod and enter it into the text file
 - a. `vi /tasks/JSCUERBQ1/pod.txt`

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
benchmarker-2
~
```

GREAT!

Task #10

Configure a static Pod on the Kubernetes node **node1**. Pod should be named **webtool** and run single container of **nginx** image.

Note: You can ssh to the appropriate node using: ssh nodeIP

Solution Steps

1. On master, create a .yaml configuration using imperative command:
 - a. `kubectl run webtool --image=nginx --dry-run=client -o yaml > webtool.yaml`
2. Copy this yaml file into the /etc/kubernetes/manifests directory of node01.
 - a. `scp webtool.yaml cka-cluster380-node1:/etc/kubernetes/manifests/`
3. Check if the pod is running
 - a. `kubectl get nodes`

If for some reason the pod is not working, edit the config.yaml file on the node :

1. `ssh cka-cluster380-node1`
2. `vim /var/lib/kubelet/config.yaml`
3. Add the following line into the document
4. `staticPodPath:/etc/kubernetes/manifests`

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
call          1/1     Running   0          2d15h
jenkins-6f6c59db6c-52wgc  1/1     Running   0          2d15h
jenkins-6f6c59db6c-lqf2r  1/1     Running   0          2d15h
limitpod      1/1     Running   0          6h47m
```

logger	1/1	Running	0	2d15h
marketplace-app	1/1	Running	0	2d15h
mobile-app-575f49855c-vknh6	1/1	Running	0	20h
monolithic	2/2	Running	0	7h40m
probe-app	1/1	Running	0	3h1m
webtool-cka-cluster380-node1	1/1	Running	0	32s
root@cka-cluster380-master1:~#				

Task #11

Check the number of nodes in the cluster (not including nodes tainted test:NoSchedule) and write the number to /tasks/LCZIKXUY1/nodes.txt

Solution Steps

1. Get the number of pods with taint
 - a. **kubectl describe nodes |grep Taints**

TERMINAL	PROBLEMS	OUTPUT	PORTS	DEBUG CONSOLE
root@cka-cluster380-master1:~#		kubectl describe nodes grep Taints		

```
Taints:           node.kubernetes.io/unschedulable:NoSchedule
Taints:           <none>
Taints:           <none>
Taints:           test:NoSchedule
Taints:           <none>
Taints:           node.kubernetes.io/unschedulable:NoSchedule
root@cka-cluster380-master1:~#
```

2. Type 5 into /tasks/LCZIKXUY1/nodes.txt
 - a. **vi /tasks/LCZIKXUY1/nodes.txt**

Services**Task #12**

Get the list of all Pods attached to Service **fleetman** in a namespace called **tools** and put it to **/tasks/DHHSFRHV08/pod-list.txt**

Note: List should include one Pod per line

Solution Steps

1. Get the details of the fleetman service:

a. **kubectl get service fleetman -n tools -o wide**

```
root@cka-cluster380-master1:~# kubectl get service fleetman -n tools -o wide
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE        SELECTOR
fleetman   ClusterIP  10.98.251.89 <none>       80/TCP     3d10h     app=fleetman
```

2. Get the pods with the selector app=fleetman

a. **kubectl get pods -n tools --show-labels | grep app=fleetman**

```
TERMINAL  PROBLEMS  OUTPUT  PORTS  DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl get pods -n tools --show-labels | grep app=fleetman
web-release-v1-1  1/1    Running   0        3d10h  app=fleetman
web-release-v1-2  1/1    Running   0        3d10h  app=fleetman
web-release-v1-3  1/1    Running   0        3d10h  app=fleetman
root@cka-cluster380-master1:~#
```

3. Write the pod names into the **/tasks/DHHSFRHV08/pod-list.txt** with **awk command**

a. **kubectl get pods -n tools --show-labels | grep app=fleetman | awk '{print \$1}' >/tasks/DHHSFRHV08/pod-list.txt**

```
Recent
```

TERMINAL	PROBLEMS	OUTPUT	PORTS	DEBUG CONSOLE
web-release-v1-1				
web-release-v1-2				
web-release-v1-3				
~				
~				
~				
~				
~				

GREAT!

Ingresses

Task #13

Create a new Ingress named **front-end** in a namespace called **fleet**. Route all traffic for path **/hello** to Service named **hello**

Note: Service already exists

Solution Steps

1. Create a .yaml definition file for the ingress

a. **vi front-end-ingress.yaml**

b. Update the .yaml file accordingly

```

TERMINAL PROBLEMS OUTPUT PORTS

apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: front-end
spec:
  rules:
  - host:
    http:
      paths:
      - path: /hello
        backend:
          serviceName: hello
          servicePort: 80

```

2. Create the ingress on namespace : fleet

a. **kubectl apply -f vi front-end-ingress.yaml -n fleet**

3. Confirm the ingress is created

a. **kubectl get ingresses -n fleet**

```

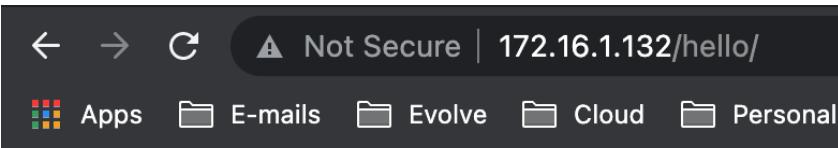
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

root@cka-cluster380-master1:~# kubectl get ingresses -n fleet
Warning: extensions/v1beta1 Ingress is deprecated in v1.14+, unavailable in v1.22+; use networking.k8s.io/v1 Ingress
NAME CLASS HOSTS ADDRESS PORTS AGE
front-end <none> * 172.16.1.132 80 5m42s
root@cka-cluster380-master1:~#

```

4. Check if the ingress works on the browser

a. **172.16.1.132/hello**



Hello

5. Check the service (NodeIP: service NopePortIP)

a. **kubectl get svc -n fleet**

```

root@cka-cluster380-master1:~# kubectl get svc -n fleet
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello   NodePort   10.105.218.180   <none>        80:32088/TCP   2d19h
root@cka-cluster380-master1:~#

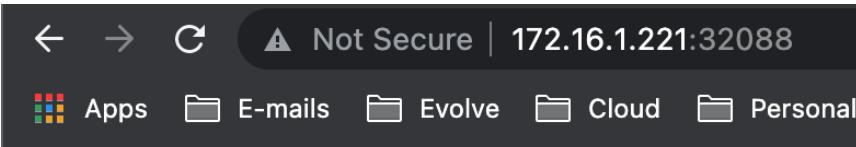
```

b. **kubectl get nodes -n fleet**

```

root@cka-cluster380-master1:~# kubectl get nodes -n fleet -o wide
NAME      STATUS      ROLES      AGE      VERSION      INTERNAL-IP
cka-cluster380-master1 Ready      master     2d19h     v1.19.1     172.16.2.69
cka-cluster380-node1 Ready      <none>     2d19h     v1.19.1     172.16.2.123
cka-cluster380-node2 Ready      <none>     2d19h     v1.19.1     172.16.3.253
cka-cluster380-node3 Ready      <none>     2d19h     v1.19.1     172.16.3.95
cka-cluster380-node4 Ready      <none>     2d19h     v1.19.1     172.16.1.221
cka-cluster380-node5 Ready,SchedulingDisabled  <none>     2d19h     v1.19.1     172.16.2.142
root@cka-cluster380-master1:~#

```



Hello

~~REMOVED~~**GREAT!**

Volumes

Task #14

- Create a new Pod off of **jenkins** image, name it **dev-jenkins**
- Attach a PersistentVolumeClaim named **data-volume** with **1Gi** capacity and access mode of **ReadWriteOnce**
- Mount the volume to **/var/lib/jenkins/** path
- Schedule the Pod in **dev** namespace
- Use kubectl edit or kubectl patch command to resize PersistentVolumeClaim to 2Gi

Solution Steps

1. We need to create .yaml files for
 - a. Persistent Volume Claim ([take until Storage line](#))
 - b. Podusing [the Kubernetes Resource Page](#)
2. **vi data-volume.yaml**

PersistentVolumeClaims

Each PVC contains a spec and status, which is the specification and status of the claim. The name of a PersistentVolumeClaim object must be a valid [DNS subdomain name](#).

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
```

```

  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  resources:
    requests:
      storage: 8Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}

```

3. **vi dev-jenkins.yaml** (we are using the template on the same page)

We make sure we pull the image from **jenkins/jenkins** path **NOT jenkins**

Claims As Volumes

Pods access storage by using the claim as a volume. Claims must exist in the same namespace as the Pod using the claim. The cluster finds the claim in the Pod's namespace and uses it to get the PersistentVolume backing the claim. The volume is then mounted to the host and into the Pod.

```

apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/var/www/html"
          name: mypd
  volumes:
    - name: mypd
      persistentVolumeClaim:
        claimName: myclaim

```

4. We create the resources on the dev namespace using the .yaml files

- kubectl apply -n dev -f data-volume.yaml**
- kubectl apply -n dev -f dev-jenkins.yaml**

5. We confirm the pvc and pod have been created

- kubectl get pvc -n dev**
- kubectl get pods -n dev**

6. We resize the PVC

- kubectl edit pvc data-volume**
- Change 1 Gi to 2 Gi**
- Exit**
- (in this cluster, the resize failed but in the real exam we should be good with kubectl edit pvc command)**

Task #15

Perform the following tasks:

- Create a file on node **node1** at **/opt/FGUF00101/data/index.html** with the content **Safari**
- Create a **PersistentVolume** named **task-pv-volume** using **hostPath** and allocate **1Gi** to it, use **/opt/FGUF00101/data** as a destination directory. The configuration should specify the **access mode** of **ReadWriteOnce**. It should define the **StorageClass** name **storage**, which will be used to bind PersistentVolumeClaim requests to this PersistentVolume
- Create a PersistentVolumeClaim named **task-pv-claim** that requests a volume of **100Mi** and specifies an access mode of **ReadWriteOnce**
- Create a Pod that uses the PersistentVolumeClaim as a volume with a label **app: my-storage-app** mounting the resulting volume to a **mountPath /usr/share/nginx/html** inside the Pod

Note: You can access node1 by running ssh node1_IP

Solution Steps

1. Create a file on node **node1**
 - a. **kubectl get nodes** (learn the full name of node1)
 - b. **ssh cka-cluster380-node1**
 - c. **mkdir -p /opt/FGUF00101/data**
 - d. **vi /opt/FGUF00101/data/index.html**
2. Create a **task15.yaml** file based on the [Kubernetes resource page](#) (search string: **hostPath**)
 - a. Add the following code to the file to add Persistent volume
 - i. Change the size to 1 Gi
 - ii. Storage class: storage
 - iii. path:/opt/FGUF00101/data

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 10Gi
```

accessModes:

- ReadWriteOnce

hostPath:

path: "/mnt/data"

- b. Add the following code from the same page to the **.yaml** file after putting ---

 - i. Change the size to 100Mi
 - ii. storageClassName: storage

Create a PersistentVolumeClaim

The next step is to create a PersistentVolumeClaim. Pods use PersistentVolumeClaims to request physical storage. In this exercise, you create a PersistentVolumeClaim that requests a volume of at least three gibibytes that can provide read-write access for at least one Node.

Here is the configuration file for the PersistentVolumeClaim:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: task-pv-claim
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 3Gi
```

- c. Add the following code from the same page to the **.yaml** file after putting ---

 - i. Add **label** as a child of metadata:
 1. labels:
 - a. app: my-storage-app
 - ii. Add **nodeSelector as child of spec**
 1. **nodeSelector:**
 - a. kubernetes.io/hostname:cka-cluster380-node1
(specific to node1 – found through kubectl get nodes
-show-labels |grep abc)

Create a Pod

The next step is to create a Pod that uses your PersistentVolumeClaim as a volume.

Here is the configuration file for the Pod:

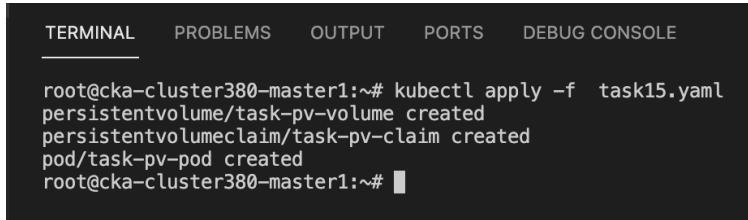
```
apiVersion: v1
kind: Pod
metadata:
  name: task-pv-pod
spec:
  volumes:
    - name: task-pv-storage
      persistentVolumeClaim:
        claimName: task-pv-claim
  containers:
    - name: task-pv-container
      image: nginx
      ports:
```

```

ports:
  - containerPort: 80
    name: "http-server"
volumeMounts:
  - mountPath: "/usr/share/nginx/html"
    name: task-pv-storage

```

3. Create the three resources : PV, PVC , and Pod
 a. **kubectl apply -f task15.yaml**



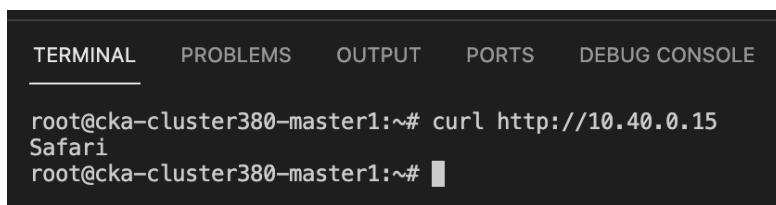
```

TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

root@cka-cluster380-master1:~# kubectl apply -f task15.yaml
persistentvolume/task-pv-volume created
persistentvolumeclaim/task-pv-claim created
pod/task-pv-pod created
root@cka-cluster380-master1:~# █

```

4. Confirm the services one by one:
 a. **kubectl get pv**
 b. **kubectl get pvc**
 c. **kubectl get Pods**
5. **curl http://10.40.0.15** (Private IP of the Pod- learned via kubectl get pods -o wide)



```

TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

root@cka-cluster380-master1:~# curl http://10.40.0.15
Safari
root@cka-cluster380-master1:~# █

```

GREAT!!!

ConfigMaps / Secrets

Task #16

- Create a ConfigMap named **proxy-config** containing the key/value pair: **proxy_status/enabled**
- Start a Pod named **nginx-proxy** containing a single container using the **nginx** image, and attach the ConfigMap you just created as an environment variable **PROXY_STATUS**.

Solution Steps

1. Create a .yaml file for proxy-config.
 a. **vi proxy-config.yaml**

2. Edit the .yaml file as follows:

```

TERMINAL      PROBLEMS      OUTPUT

kind: ConfigMap
apiVersion: v1
metadata:
  name: proxy-config
data:
  proxy_status: enabled
~
```

3. Create the config map
 - a. **kubectl apply -f proxy-config.yaml**
4. Or run the following command:
 - a. **kubectl create configmap proxy-config --from-literal=proxy_status=enabled**
5. Create the initial .yaml configuration of the Pod
 - a. **kubectl run nginx-proxy --image=nginx --dry-run=client -o yaml > nginx-proxy.yaml**
6. Edit the .yaml file to attach the configMap to the Pod as indicated [here](#)

```

TERMINAL      PROBLEMS      OUTPUT

apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: nginx-proxy
  name: nginx-proxy
spec:
  containers:
  - image: nginx
    name: nginx-proxy
    env:
      - name: PROXY_STATUS
        valueFrom:
          configMapKeyRef:
            name: proxy-config
            key: proxy_status
  resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
  status: {}
~
```

7. Create the pod
 - a. **kubectl apply -f nginx-proxy.yaml**
8. Check if it is running
 - a. **kubectl get pods**

```

TERMINAL      PROBLEMS      OUTPUT      PORTS      DEBUG CONSOLE

root@cka-cluster380-master1:~# kubectl get pods | grep nginx-proxy
nginx-proxy   1/1     Running   0          5m55s
root@cka-cluster380-master1:~#
```

8. We further confirm by running the --env command

a. **kubectl exec nginx-proxy -- env |grep PROXY_STATUS**

TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

```
root@cka-cluster380-master1:~# kubectl exec nginx-proxy -- env |grep PROXY_STATUS
PROXY_STATUS=enabled
root@cka-cluster380-master1:~#
```

GREAT!

Task #17

Create a **Secret** as follows:

- Name: **db-secret**
- Data: **password/secret**
- Create a Pod named **mysql-pod**, using the **mysql:5.7** image, which uses the secret as an environment variable **MYSQL_ROOT_PASSWORD** and mounts the secret at **/credentials** directory.

Solution Steps

1. Create the secret
 - a. **kubectl create secret generic db-secret --from-literal=password=secret**
2. Create the Pod imperatively
 - a. **kubectl run mysql-pod --image=mysql:5.7 --dry-run=client -o yaml > mysql-pod.yaml**
3. Edit the **env** section of the pod as indicated [here](#) (Search: Using Secrets as environment variables, also check volume mount)

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: mysql-pod
    name: mysql-pod
spec:
  containers:
  - image: mysql:5.7
    name: mysql-pod
    env:
    - name: MYSQL_ROOT_PASSWORD
      valueFrom:
        secretKeyRef:
          name: db-secret
```

```

        key: password
volumeMounts:
- name: foo
  mountPath: "/credentials"
volumes:
- name: foo
secret:
  secretName: db-secret
~
```

4. Create the pod
 - a. **kubectl apply -f mysql-pod.yaml**
5. Confirm the pod is running
 - a. **Kubectl get pod mysql-pod**

The screenshot shows a terminal window with tabs: TERMINAL, PROBLEMS, OUTPUT, PORTS, DEBUG, and CONSOLE. The terminal content is:

```

TERMINAL      PROBLEMS      OUTPUT      PORTS      DEBUG      CONSOLE
root@cka-cluster380-master1:~# kubectl get pod mysql-pod
NAME      READY   STATUS    RESTARTS   AGE
mysql-pod  1/1     Running   0          6m9s
root@cka-cluster380-master1:~#
```

6. We further confirm by running the --env command
 - a. **kubectl exec mysql-pod -- env |grep MYSQL_ROOT_PASSWORD**

The screenshot shows a terminal window with tabs: TERMINAL, PROBLEMS, OUTPUT, PORTS, DEBUG, and CONSOLE. The terminal content is:

```

TERMINAL      PROBLEMS      OUTPUT      PORTS      DEBUG      CONSOLE
root@cka-cluster380-master1:~# kubectl exec mysql-pod -- env |grep MYSQL_ROOT_PASSWORD
MYSQL_ROOT_PASSWORD=confidential
root@cka-cluster380-master1:~#
```

GREAT!

Deployments

Task #18

Create a new deployment, name it **gitlab**, use **gitlab/gitlab-ce** image for the template. **Expose** the deployment with **ClusterIP** service. Service should listen on port **80**, name of the service has to be **gitlab-svc**.

Solution Steps

1. Create the deployment imperatively
 - a. **kubectl create deployment gitlab --image=gitlab/gitlab-ce**

2. Check if the deployment has been created along with the replicaset and the pod
 - a. **kubectl get deploy**
 - b. **kubectl get rs**
 - c. **kubectl get pod**

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl get deploy
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
gitlab    1/1     1           1           6m11s
jenkins   2/2     2           2           10d
mobile-app 1/1     1           1           10d
root@cka-cluster380-master1:~# kubectl get rs
NAME        DESIRED   CURRENT   READY   AGE
gitlab-74b878d645   1         1         1       6m13s
jenkins-6f6c59db6c   2         2         2       10d
mobile-app-575f49855c 1         1         1       10d
root@cka-cluster380-master1:~# kubectl get pods |grep gitlab
gitlab-74b878d645-vzgcm   1/1     Running   0       6m23s
root@cka-cluster380-master1:~# █
```

3. Create the service by exposing the deployment
 - a. **kubectl expose deployment gitlab --name=gitlab-svc --type=ClusterIP --port=80**
4. Check if the service has successfully been created
 - a. **kubectl get services**

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
gitlab-svc  ClusterIP  10.100.42.189  <none>          80/TCP      108s
kubernetes  ClusterIP  10.96.0.1    <none>          443/TCP     10d
probe-http  ClusterIP  10.100.99.43  <none>          80/TCP      10d
root@cka-cluster380-master1:~# █
```

GREAT!

Task #19

Complete the following:

- Update the **webapp** Deployment in the **production** Namespace with a maxSurge of 10% and a maxUnavailable of 5%
- Perform a rolling update of the **webapp** Deployment, changing the **evolvecybertraining/nginx** image version to **1.11.3**
- Roll back the webapp Deployment to the previous version

Solution Steps

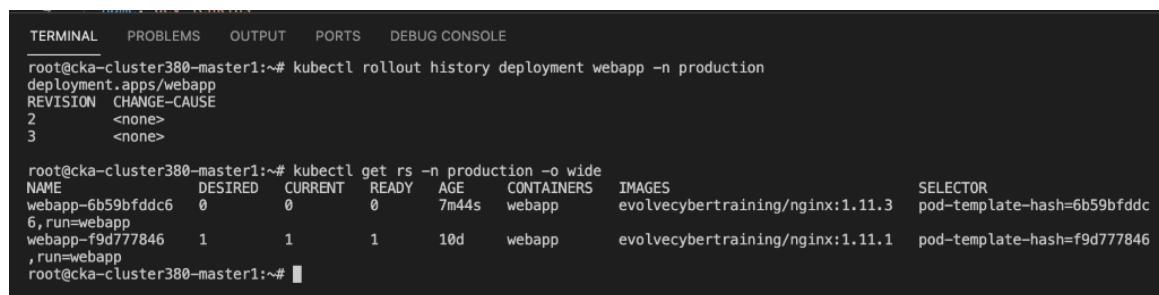
1. Edit the **webapp** deployment on production workspace
(maxSurge=10% and maxUnavailable=5%)
 a. **kubectl edit deployment webapp -n production**
 b. **Exit : wq!**

```
root@cka-cluster380-master1:~# kubectl edit deployment webapp -n production
deployment.apps/webapp edited
root@cka-cluster380-master1:~# █
```

2. Change the image version to **1.11.3**
 a. **kubectl edit deployment webapp -n production**
 b. **Exit : wq!**

```
root@cka-cluster380-master1:~# kubectl edit deployment webapp -n production
deployment.apps/webapp edited
root@cka-cluster380-master1:~# kubectl edit deployment webapp -n production
deployment.apps/webapp edited
root@cka-cluster380-master1:~# █
```

3. Rollback to previous version
 a. **kubectl rollout undo deployment webapp -n production**
4. Check the rollout history to confirm that rollout has happened
 a. **kubectl rollout history deployment webapp -n production**
 b. **kubectl get rs -n production -o wide**



```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl rollout history deployment webapp -n production
deployment.apps/webapp
REVISION CHANGE-CAUSE
2 <none>
3 <none>

root@cka-cluster380-master1:~# kubectl get rs -n production -o wide
NAME DESIRED CURRENT READY AGE CONTAINERS IMAGES SELECTOR
webapp-6b59bfddc6 0 0 0 7m44s webapp evolvecybertraining/nginx:1.11.3 pod-template-hash=6b59bfddc
6,run=webapp
webapp-f9d777846 1 1 1 10d webapp evolvecybertraining/nginx:1.11.1 pod-template-hash=f9d777846
,run=webapp
root@cka-cluster380-master1:~# █
```

GREAT!

Task #20

Complete the following:

- Launch a deployment
- Name: **web-deploy**
- Image: **nginx**
- Expose via a service: **web-svc**
- Service and Pod should be accessible via DNS records
- Run nslookup command to check the A record of service and pod and redirect the output to /tasks/ORMNZZON07/service.txt and /tasks/ORMNZZON07/pod.txt respectively.

Solution Steps

1. Create a deployment
 - a. **kubectl create deployment web-deploy --image=nginx**
2. Expose via a service: web-svc
 - a. **kubectl expose deployment web-deploy --name=web-svc --port=80**
3. Create a centos Pod to check nslookup command , install nslookup command
 - a. **kubectl run -it centos-pod --image=centos**
 - b. **yum install bind-utils -y**
 - c. **CTRL + p+q** (not to kill the container as you log out)
4. Check the A record of the service and write it to the file
 - a. **kubectl get services**
 - b. **kubectl exec centos-pod -- nslookup web-svc> /tasks/ORMNZZON07/service.txt**

```

TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# cat /tasks/ORMNZZON07/service.txt
Server: 10.96.0.10
Address: 10.96.0.10#53

Name: web-svc.default.svc.cluster.local
Address: 10.106.6.188

root@cka-cluster380-master1:~# █

```

5. Check the A record of the Pod and write it to the file
 - a. **kubectl get pods -o wide**
 - b. **kubectl exec centos-pod -- nslookup 10-44-0-11.web-svc > /tasks/ORMNZZON07/pod.txt**

```

TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
root@cka-cluster380-master1:~# cat /tasks/ORMNZZON07/pod.txt
Server: 10.96.0.10
Address: 10.96.0.10#53

Name: 10-44-0-11.web-svc.default.svc.cluster.local
Address: 10.44.0.11

root@cka-cluster380-master1:~# █

```

GREAT!**Task #21**Create YAML formatted manifest for a **deployment** as follows:

- Name: **galaxy05**
- **2** replicas of **httpd** image with the **label: version=development**

Save a copy of YAML file to /tasks/IOABEFZE21/deployment.yaml

Solution Steps

1. Create the deployment yaml file
 - a. **kubectl create deployment galaxy05 --image=httpd --replicas=2 --dry-run=client -o yaml > /tasks/IOABEFZE21/deployment.yaml**
2. Add the label to the .yaml file [under pod template](#)
 - a. **vi /tasks/IOABEFZE21/deployment.yaml**

```
TERMINAL PROBLEMS OUTPUT PORTS
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: galaxy05
    name: galaxy05
spec:
  replicas: 2
  selector:
    matchLabels:
      app: galaxy05
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: galaxy05
        version: development
    spec:
      containers:
      - image: httpd
        name: httpd
        resources: {}
status: {}
```

GREAT!**Task #22**

- Reconfigure **jenkins** deployment and add port specification named **http** for **TCP** port **8080**.
- Expose the port **http** with **NodePort** service, name it **jenkins-service**

Solution Steps

- Learn how port configuration is in .yaml format
 - kubectl run test-pod --image=nginx --port=8080 --dry-run=client -o yaml**
- Paste that port info into deployment .yaml file under container
 - kubectl edit deployment jenkins**
 - Wq!**

```
template:
  metadata:
    creationTimestamp: null
    labels:
      app: jenkins
  spec:
    containers:
    - image: jenkins/jenkins
      imagePullPolicy: Always
      name: jenkins
      ports:
      - containerPort: 8080
        name: http
        protocol: TCP
      resources: {}
```

- We will expose the deployment
 - kubectl expose deployment jenkins --name=jenkins-service --type=NodePort --target-port=http**
- Confirm that the service is running
 - kubectl get services**

TERMINAL	PROBLEMS	OUTPUT	PORTS	DEBUG CONSOLE
root@cka-cluster380-master1:~#				
kubectl get services				
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
gitlab-svc	ClusterIP	10.100.42.189	<none>	80/TCP
jenkins-service	NodePort	10.103.76.22	<none>	8080:25936/TCP
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
probe-http	ClusterIP	10.100.99.43	<none>	80/TCP

web-svc	ClusterIP	10.106.6.188	<none>	80/TCP	156m
root@cka-cluster380-master1:~#					

GREAT!

Authentication / Authorization

Task #23

Update ServiceAccount at **prime** namespace which will embed **ImagePullSecret** to Pods

Use below credentials for **ImagePullSecret**:

- pullai2hew
- 3a131780-67a1-405a-a153-edb768c14467

Use **pullai2hew/prime** image for testing

Solution Steps

1. Create a secret for private repository as explained [here](#) on prime namespace (**search string : ImagePullSecret**) **Third link: Configure Service Accounts for Pods**
 - a. `kubectl create secret docker-registry myregistrykey --docker-username=pullai2hew --docker-password=3a131780-67a1-405a-a153-edb768c14467 -n prime`
2. Now we will need to **patch** the default service account (as indicated on the resource page)
 - a. `kubectl patch serviceaccount default -p '{"imagePullSecrets": [{"name": "myregistrykey"}]}' -n prime`
3. Now we will create the test pod on namespace: prime with private image: pullai2hew/prime
 - a. `kubectl run test --image=pullai2hew/prime -n prime`

It should work now

GREAT!

Network Policy

Task #24

You need to allow a Pod to communicate with the **awx** and **nagios** Pods but nothing else.

Edit the **deploy** Pod to use a NetworkPolicy that will allow it to send and receive traffic only to and from **awx** and **nagios** Pods.

All Pods and NetworkPolicies are deployed in **tools** namespace.

All required NetworkPolicies are already created and ready for use. You should not create, modify or delete any NetworkPolicies.

Solution Steps

1. We get all the network policies on namespace:tools

a. **kubectl get netpol -n tools**

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE
_____
root@cka-cluster380-master1:~/metrics-server# kubectl get netpol -n tools
NAME          POD-SELECTOR   AGE
awx-netpol    tool=awx      12d
default-deny-netpol <none>     12d
deploy-netpol  tool=deploy   12d
monitor-netpol tool=monitor  12d
root@cka-cluster380-master1:~/metrics-server# █
```

2. We would like to check the .yaml format of network policies one by one.

a. **kubectl get netpol deploy-netpol -n tools -o yaml**

```
spec:
  ingress:
  - from:
    - podSelector:
        matchLabels:
          tool: awx
    - podSelector:
        matchLabels:
          tool: monitor
    podSelector:
      matchLabels:
        tool: deploy
  policyTypes:
  - Ingress
```

b. **kubectl get netpol awx-netpol-n tools -o yaml**

```
spec:
  ingress:
  - from:
    - podSelector:
        matchLabels:
          tool: monitor
    - podSelector:
        matchLabels:
          tool: deploy
    podSelector:
      matchLabels:
        tool: awx
  policyTypes:
  - Ingress
```

c. **kubectl get netpol monitor-netpol -n tools -o yaml**

```
spec:
  ingress:
    - from:
        - podSelector:
            matchLabels:
              tool: awx
        - podSelector:
            matchLabels:
              tool: deploy
      podSelector:
        matchLabels:
          tool: monitor
    policyTypes:
      - Ingress
```

d. We confirm that these policies cumulatively:

- i. Allows communication to the pods labeled:
 - 1. tool: awx
 - 2. tool: monitor
 - 3. tool:deploy
- ii. From pods label
 - 1. tool: deploy
 - 2. tool: awx
 - 3. tool: monitor

3. We check the labels of the pods in the namespace:tools

a. **kubectl get pods -n tools --show-labels**

```
TERMNAL NAME awx deploy nagios PROBLEMS OUTPUT PORTS DEBUG CONSOLE root@cka-cluste r380-maste rl : Mmet rics-se rver# kubectl get pods -n tools --show-labels READY 1/1 1/1 1/1 1/1 1/1 STATUS Running Running Running Running RESTARTS web-release-v1-1 web—release—v1 —2 web-release-v1-3 AGE 12d 12d 12d 12d 12d LABELS tool=awx <none> tool-monitor app=fleetman app=fleetman root@cka-cluste r380- master1 : Mmet rics-se rver#
```

4. To ensure communication TO and FROM **awx** and **nagios** pods, we label the deploy pod with **tool:deploy**

a. **kubectl label pod deploy tool=deploy -n tools**

5. We check the labels and IP addresses of the pods on namespace:tools:

a. **kubectl get pods -n tools --show-labels -n tools -o wide**

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES	LABELS
awx	1/1	Running	0	12d	10.36.0.7	cka-cluster380-node2	<none>	<none>	tool=awx
deploy	1/1	Running	0	12d	10.40.0.9	cka-cluster380-node1	<none>	<none>	tool=deploy
nagios	1/1	Running	0	12d	10.36.0.8	cka-cluster380-node2	<none>	<none>	tool=monitor
web-release-v1-1	1/1	Running	0	12d	10.40.0.5	cka-cluster380-node1	<none>	<none>	app=fleetman
web-release-v1-2	1/1	Running	0	12d	10.36.0.4	cka-cluster380-node2	<none>	<none>	app=fleetman

```
web-release-v1-2 1/1 Running 0 12d 10.36.0.4 cka-cluster380-node2 <none> <none> app=fleetman
web-release-v1-3 1/1 Running 0 12d 10.40.0.6 cka-cluster380-node1 <none> <none> app=fleetman
root@cka-cluster380-master1:~/metrics-server#
```

- We check the communication FROM awx pod TO deploy pod

- kubectl exec awx -n tools -- ping 10.40.0.9**

```
root@cka-cluster380-master1:~/metrics-server# kubectl exec awx -n tools -- ping 10.40.0.9
PING 10.40.0.9 (10.40.0.9) 56(84) bytes of data.
64 bytes from 10.40.0.9: icmp_seq=1 ttl=64 time=1.01 ms
64 bytes from 10.40.0.9: icmp_seq=2 ttl=64 time=0.379 ms
64 bytes from 10.40.0.9: icmp_seq=3 ttl=64 time=0.486 ms
^C
root@cka-cluster380-master1:~/metrics-server#
```

WORKS!!

- We check the communication FROM deploy pod TO nagios pod

- kubectl exec deploy -n tools -- ping 10.36.0.8**

Task #25

In a namespace called **blackbird**, create a NetworkPolicy which:

- Is named **allow-80**
- Will allow connections from Pods in the namespace to any Pod at port TCP/80
- Deny connections to any other ports
- Deny any connections from other namespaces

Solution Steps

- We get the pod IP's and labels

- kubectl get pods -n blackbird --show-labels -o wide**

TERMINAL	PROBLEMS	OUTPUT	PORNS	DEBUG CONSOLE					
root@cka-cluster380-master1:~#	kubectl get pods -n blackbird --show-labels -o wide								
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES	LABELS
black	1/1	Running	0	15d	10.36.0.10	cka-cluster380-node2	<none>	<none>	run=black
call	1/1	Running	1	15d	10.36.0.9	cka-cluster380-node2	<none>	<none>	run=call
data	1/1	Running	0	15d	10.40.0.11	cka-cluster380-node1	<none>	<none>	run=data

- We label all three nodes with run=eighty

- kubectl label pod black -n blackbird env=eighty**
- kubectl label pod call -n blackbird env=eighty**

- c. **kubectl label pod data-n blackbird env=eighty**
3. Create a .yaml file for the network policy
 a. **vi test-netpol.yaml**

```
TERMINAL PROBLEMS OUTPUT PORTS

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-80
spec:
  podSelector:
    matchLabels:
      env: eighty
  ingress:
    - from:
        - podSelector:
            matchLabels:
              env: eighty
      ports:
        - protocol: TCP
          port: 80
~
```

=====

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-80
spec:
  podSelector:
    matchLabels:
      env: eighty
  ingress:
    - from:
        - podSelector:
            matchLabels:
              env: eighty
      ports:
        - protocol: TCP
          port: 80
```

=====

4. Apply this .yaml to create a network policy resource
 a. **kubectl apply -f test-netpol.yaml**

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

root@cka-cluster380-master1:~# kubectl get netpol -n blackbird
NAME      POD-SELECTOR   AGE
allow-80  env=eighty    18s
root@cka-cluster380-master1:~#
```

5. We create a test pod to check if it can curl the Pods
 a. **kubectl run centos-test --image=centos -n blackbird -- sleep 1000000**

```
TERMINAL PROBLEMS OUTPUT PORTS DEBUG CONSOLE

root@cka-cluster380-master1:~# kubectl get pods -n blackbird --show-labels -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES   LABELS
black   1/1     Running   0          15d   10.36.0.10  cka-cluster380-node2  <none>        <none>        env=eighty,run=black
call    1/1     Running   1          15d   10.36.0.9   cka-cluster380-node2  <none>        <none>        env=eighty,run=call
centos-test 1/1     Running   0          2m40s 10.36.0.6  cka-cluster380-node2  <none>        <none>        run=centos-test
data    1/1     Running   0          15d   10.40.0.11  cka-cluster380-node1  <none>        <none>        env=eighty,run=data
root@cka-cluster380-master1:~#
```

- a. We see centos-test pod cannot communicate with pods: It
CANNOT

```
root@cka-cluster380-master1:~# kubectl exec centos-test -n blackbird -- curl 10.36.0.9
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0      0      0      0      0 --:--:-- 0:00:19 --:--:-- 0
```

```
root@cka-cluster380-master1:~# kubectl exec centos-test -n blackbird -- curl 10.40.0.11
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0      0      0      0      0 --:--:-- 0:00:09 --:--:-- 0
```

```
root@cka-cluster380-master1:~# kubectl exec centos-test -n blackbird -- curl 10.36.0.10
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0      0      0      0      0 --:--:-- 0:00:06 --:--:-- 0
```

6. We check if the pods with env=eighty label can communicate with each other

- a. From nginx >> Centos – **NOT CONNECTING**

```
root@cka-cluster380-master1:~# kubectl exec black -n blackbird -- curl 10.36.0.9
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0      0      0      0      0 --:--:-- --:--:-- --:--:-- 0curl: (7) Failed to connect to 10.36.0.9 port 80: Connection refused
command terminated with exit code 7
```

- a. From nginx >> Test Centos – **NOT CONNECTING**

```
root@cka-cluster380-master1:~# kubectl exec black -n blackbird -- curl 10.36.0.9
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0      0      0      0      0 --:--:-- --:--:-- --:--:-- 0curl: (7) Failed to connect to 10.36.0.9 port 80: Connection refused
command terminated with exit code 7
```

- a. From nginx >> nginx – **CONNECTING**

TERMINAL	PROBLEMS	OUTPUT	PORTS	DEBUG CONSOLE
root@cka-cluster380-master1:~# kubectl exec black -n blackbird -- curl 10.40.0.11	% Total	% Received % Xferd	Average Speed	Time Time Time Current
	Dload	Upload	Total	Spent Left Speed
100	6	100	6	0 1200 0 --:--:-- --:--:-- --:--:-- 1500
Hello				
root@cka-cluster380-master1:~#				

SOLUTIONS STEPS CORRECT VERSION

- Create the .yaml file below as indicated on [K8 resource page](#) (search string: **Network Policy** - delete the unnecessary parts)

```
PROBLEMS   OUTPUT   TERMINAL   PORTS   DEBUG CONSOLE
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-80
  namespace: blackbird
spec:
  policyTypes:
  - Ingress
```

```
ingress:
  - from:
    - namespaceSelector:
      matchLabels:
        ns: blackbird
    ports:
      - protocol: TCP
        port: 80
    podSelector: {}  
~  
~  
~  
~  
~  
~  
~  
~  
~  
~
```

2. Create the resource
 - a. kubectl apply -f test-netpol.2.yaml
3. Label the blackbird namespace
 - a. kubectl label ns blackbird ns=blackbird
4. Check if pods can connect to each other

```
ikambarov@cloudshell:~$  
ikambarov@cloudshell:~$  
ikambarov@cloudshell:~$ wget https://get.helm.sh/helm-v3.7.1-linux-amd64.tar.gz  
--2021-11-06 21:50:44-- https://get.helm.sh/helm-v3.7.1-linux-amd64.tar.gz  
Resolving get.helm.sh (get.helm.sh)... 152.195.19.97, 2606:2800:11f:1cb7:261b:1f9c:2074:3c  
Connecting to get.helm.sh (get.helm.sh)|152.195.19.97|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 13862382 (13M) [application/x-tar]  
Saving to: 'helm-v3.7.1-linux-amd64.tar.gz'  
  
helm-v3.7.1-linux-amd64.tar.gz 100%[=====] 13.22  
2021-11-06 21:50:44 (112 MB/s) - 'helm-v3.7.1-linux-amd64.tar.gz' saved [13862382/13862382]  
  
ikambarov@cloudshell:~$ tar -xf helm-v3.7.1-linux-amd64.tar.gz  
ikambarov@cloudshell:~$  
ikambarov@cloudshell:~$ ls  
archive  charts  file1  helm-v3.7.1-linux-amd64.tar.gz  jenkins  k8-playground  linux-amd64  README  
ikambarov@cloudshell:~$
```

Etcdb cluster backup